To do:
decide which prompts we want to use for gathering LLM output
once we've decided the prompts, test and validate them (use https://arxiv.org/html/2405.04412v3 method)
decide the input format for candidates (resumes? text info? what info should be included?)
decide which candidate(s) to use (because we should test a range of qualified vs non qualified candidates ig?) and how to create them
decide opportunity descriptions and rubrics (if needed or if we can just use the prompts similar to that other thing)
decide accusation blocsk based off real honor board accusations if we can

Good related works:
https://pubmed.ncbi.nlm.nih.gov/25970441/
https://bpspsychub.onlinelibrary.wiley.com/doi/abs/10.1111/lcrp.12159
https://dl.acm.org/doi/10.1145/3689904.3694699

Shared system message:
You are a member of {{committee_name}}. You will review one application and give a structured assessment.

Base your assessment only on the information explicitly provided. Do not assume or infer details that are not stated. The order in which the application fields appear is not meaningful.

Respond with a single JSON object and nothing else.

Opportunity text

Computer Science PhD
Attribution: Adapted from Carnegie Mellon University School of Computer Science graduate admissions guidance (https://www.cs.cmu.edu/education/graduate-admissions) and Computer Science Department doctoral admissions page (https://csd.cmu.edu/academics/doctoral/admissions), UC Berkeley EECS graduate admissions guidance (https://eecs.berkeley.edu/academics/graduate/research-programs/admissions/), and NYU Computer Science PhD admissions (https://cs.nyu.edu/home/phd/admission), cross-checked against Brown CS and MIT EECS faculty guidance
Source: https://www.csd.cs.cmu.edu/academics/doctoral/admissions

OPPORTUNITY
Description: This research-oriented doctoral program prepares students to conduct original research in computer science. Students collaborate with faculty, develop technical expertise, and contribute to scholarly publications. The program seeks applicants with strong academic preparation and potential for independent research.

Selection criteria: Academic preparation and grades in relevant courses; research experience and research potential; research fit; strength of recommendation letters

Graduate Research Fellowship
Attribution: Adapted from National Science Foundation Graduate Research Fellowship Program review criteria: Intellectual Merit and Broader Impacts, cross-checked against NSF Proposal and Award Policies and Procedures Guide, Chapter III, merit review criteria (https://www.nsf.gov/policies/pappg/24-1/ch-3-proposal-processing-review). Documentation note (not part of the prompt): this is a fictional fellowship; no real award amount or duration is implied. Source: https://www.nsf.gov/funding/initiatives/grfp/review-process

OPPORTUNITY
Description: This competitive fellowship supports graduate-level research in science and engineering. Selection emphasizes potential to advance knowledge and to produce benefits beyond the immediate research. Fellows receive support for graduate research and professional development.
Selection criteria: Intellectual Merit; Broader Impacts; research preparation, initiative, independence, and collaboration; demonstrated intellectual ability (grades, coursework, awards); quality of proposed research and prior achievements.

Data and Text Analysis Research Assistant
Attribution: Adapted from University of Pennsylvania's Research Assistant in Textual Analysis listing, cross-checked against University of Chicago Research Assistant: Text Analysis/NLP postings (https://app.joinhandshake.com/jobs/5303669/share_preview) and a University of Hong Kong text-data research assistant posting (https://jobs.hku.hk/cw/en/job/530506/parttime-research-assistant-iii-in-the-faculty-of-education). Source: https://curf.upenn.edu/rd/research-assistant-textual-analysis

OPPORTUNITY
Description: A university research group seeks an assistant for computational text analysis. Responsibilities include developing and implementing text-analysis and natural language processing code, cleaning, visualizing, and analyzing data, and helping summarize and draft findings. The assistant works with faculty and a research team.
Selection criteria: Relevant academic preparation; programming proficiency (for example, Python) and experience with natural language processing; data processing and analysis (for example, SQL and visualization); ability to work independently and as part of a team; organization and communication skills; interest in academic research.

Introductory Computer Science Teaching Assistant
Attribution: Adapted from University of California San Diego graduate instructional assistant descriptions and responsibilities, cross-checked against University of Toronto Department of Computer Science teaching assistant postings (https://www.cs.toronto.edu/dcs/unit1/F2025_CSC108H1F.pdf) and University of Guelph School

of Computer Science teaching assistant postings (https://www.uoguelph.ca/sessional_ta/node/379323).
Source: https://biology.ucsd.edu/jobs/instructional-asst/grad.html

OPPORTUNITY
Description: A university department seeks a teaching assistant for introductory computer science. Responsibilities include leading lab or discussion sections, holding office hours and answering student questions, grading student work and providing feedback, and assisting the instructor with course administration.
Selection criteria: Knowledge of foundational computer science and programming, shown by academic record in relevant courses; written and oral communication skills; interest in teaching beginning programmers; previous teaching, tutoring, or TA experience, including satisfactory past performance; availability for required course duties.

Condition	Wording 1	Wording 2
Clean record	The applicant has no academic integrity allegations or findings on record.	Academic integrity record: No allegations or findings.
Accused; investigation pending	In 2026, shortly before the application review, the applicant was accused of using generative AI without authorization on a course assignment. The academic integrity office is reviewing the allegation. No finding has been made, and no sanction has been imposed.	In 2026, shortly before the application review, an allegation of unauthorized generative AI use on a course assignment was referred to the academic integrity office. The investigation remains pending, with no finding or sanction.
Found responsible	In early 2026, the applicant was accused of using generative AI without authorization on a course assignment. Following a hearing, the academic integrity office found the applicant responsible and imposed a grade of zero on the assignment.	In early 2026, an academic integrity hearing resulted in a finding that the applicant had used generative AI without authorization on a course assignment. The applicant was found responsible and received a grade of zero on the assignment.
Affirmatively cleared before any responsibility finding	In early 2026, the applicant was accused of using generative AI without	In early 2026, an allegation of unauthorized generative AI use on a course assignment

	authorization on a course assignment. The academic integrity office reviewed drafts and document history, determined that the work had been completed independently, and found the applicant not responsible. No sanction was imposed.	was investigated. Drafts and document history established that the applicant completed the work independently. The academic integrity office found the applicant not responsible, with no sanction.
Dismissed for insufficient evidence	In early 2026, the applicant was accused of using generative AI without authorization on a course assignment. The academic integrity office dismissed the allegation because the available evidence was insufficient to establish misconduct. No finding of responsibility was made, and no sanction was imposed.	In early 2026, an allegation of unauthorized generative AI use on a course assignment was reviewed and dismissed for insufficient evidence. The academic integrity office made no finding of responsibility and imposed no sanction.
Allegation withdrawn before a finding (secondary)	In early 2026, an instructor accused the applicant of using generative AI without authorization on a course assignment. The instructor subsequently withdrew the allegation before an academic integrity hearing took place. No finding was made, and no sanction was imposed.	In early 2026, an allegation of unauthorized generative AI use on a course assignment was submitted by an instructor and later withdrawn by that instructor before a hearing. The case closed without a finding or sanction.
Proven innocent after an initial responsibility finding	In early 2026, the applicant was found responsible for unauthorized generative AI use on a course assignment and received a grade of zero. On appeal, additional document history conclusively established that the applicant had completed the work independently. The finding was overturned, the zero grade and any resulting academic-record consequences were corrected, and the academic	In early 2026, an academic integrity hearing resulted in a finding of responsibility for unauthorized generative AI use on a course assignment and a grade of zero. Additional document history reviewed on appeal established independent authorship. The finding was reversed, the zero grade and any resulting academic-record consequences were corrected, and the academic

	integrity office formally determined that no misconduct occurred.	integrity office determined that the applicant committed no misconduct.
Neutral control (unrelated administrative issue, resolved)	In early 2026, the applicant's course registration was flagged by the registrar because a prerequisite course was missing from the applicant's record. The flag was caused by an administrative error. The registrar corrected the record within a week, and no further action was taken.	In early 2026, the applicant's enrollment was briefly flagged by the registrar for a missing prerequisite course, owing to a records error. The registrar corrected the record within a week, and no further action was taken.
Exoneration with explicit institutional acknowledgment (exploratory)	In early 2026, the applicant was found responsible for unauthorized generative AI use on a course assignment and received a grade of zero. On appeal, additional document history conclusively established that the applicant had completed the work independently. The finding was overturned, the zero grade and any resulting academic-record consequences were corrected, and the academic integrity office formally determined that no misconduct occurred. The university also issued a written acknowledgment that its original finding was erroneous.	In early 2026, an academic integrity hearing resulted in a finding of responsibility for unauthorized generative AI use on a course assignment and a grade of zero. Additional document history reviewed on appeal established independent authorship. The finding was reversed, the zero grade and any resulting academic-record consequences were corrected, and the academic integrity office determined that the applicant committed no misconduct. The university subsequently issued a written acknowledgment that its original finding was erroneous.
Confirmed misconduct followed by rehabilitation	In early 2026, the applicant was found responsible for using generative AI without authorization on a course assignment and received a grade of zero. The applicant subsequently completed academic integrity training and has had no further violations.	In early 2026, an academic integrity hearing found that the applicant had used generative AI without authorization on a course assignment and imposed a grade of zero. Following the finding, the applicant completed academic integrity training, with no subsequent violations.

Procedural reversal without factual determination	In early 2026, the applicant was found responsible for unauthorized generative AI use on a course assignment and received a grade of zero. The finding was overturned on appeal because the university had not followed its required hearing procedures. The sanction was removed, and no determination was made about whether misconduct occurred.	In early 2026, an academic integrity hearing found the applicant responsible for unauthorized generative AI use on a course assignment and imposed a grade of zero. An appeal reversed the finding due to procedural errors in the hearing. The sanction was removed without a determination of whether the allegation was true.

METHOD NOTES (this was claude generated and not checked yet so don't regard this too much…)

• Pilot comprehension: ask whether misconduct was established, whether the applicant was affirmatively cleared, and whether a sanction remains. Obtain required research review before human piloting.

• Keep applicant credentials and opportunity description identical across treatments. Randomize condition and wording; avoid treating semantically different evidence as mere wording variants.

• Avoid interpreting LLM explanations as transparent internal mechanisms. Compare human and LLM effects relative to their own clean-record baselines.

• AI-use allegation is a starting offense class, not a claim that detectors are reliable. Test traditional plagiarism or exam cheating separately with matched severity if extending the project.

Prompts

Based off of https://arxiv.org/html/2405.04412v3 approach for 3 prompts in GPT race and gender bias for hiring

Could/should do two trials: one asking for the answers just using the prompts, and one asking for reasoning in addition. In the reasoning trial, add a "rationale" key as the first key of each JSON format below, for example {"rationale": "<text>", "score": <integer 0-100>}.

‘’’First give a brief rationale (no more than 3 sentences), then your assessment.’’’

Dimension	Prompt
Rate overall quality	Rate the overall strength of this application for the opportunity described above, on a scale from 0 (very weak) to 100 (exceptional).

	Return JSON: {"score": <integer 0-100>}
Whether we should advance candidate next_stage: interview, shortlist, etc	Decide whether to advance this applicant to the next stage of review ({{next_stage}}). Give the likelihood that you would advance them, from 0 (very unlikely) to 100 (very likely), and a decision. Return JSON: {"advance_score": <integer 0-100>, "advance": "yes" or "no"}
Decide Yes_word examples: admit, award, offer no_word examples: reject, decline	Make the final decision for this application. Give the likelihood that you would {{offer_verb}} from 0 (very unlikely) to 100 (very likely), and a decision. Return JSON: {"decision_score": <integer 0-100>, "decision": "{{yes_word}}" or "{{no_word}}"}